

# **Storage and Indexing**



# Memory Types

---

- **DRAM**
  - Volume. Up to 256 Gb
  - Cost. ~2-3\$ per Gb
  - Access. 1-10 ns
- **SSD**
  - Volume. Up to 15 Tb
  - Cost. ~0.1\$ per Gb
  - Access. 0.1-0.2 ms
- **HDD**
  - Volume. Up to 20 Tb
  - Cost. ~0.03\$ per Gb
  - Access. 5-100 ms

# Hard drives specificity

---



CITY UNIVERSITY  
LONDON

- Large search time
- Speed of reads
  - Consecutive - reasonable
  - Random - slow
- We need to reduce the number of accesses
  - Hopefully, making them consecutive

# Pages

---



CITY UNIVERSITY  
LONDON

- The memory of hard drives is split into pages
- Read loads into cache
- Processing is faster
- Write is also slow
- It is better to store data in consecutive pages

# Record load

---



CITY UNIVERSITY  
LONDON

- We ask for a record
- Record Manager gives us the corresponding page
- Disk Manager loads the corresponding page

# Data Storage

---



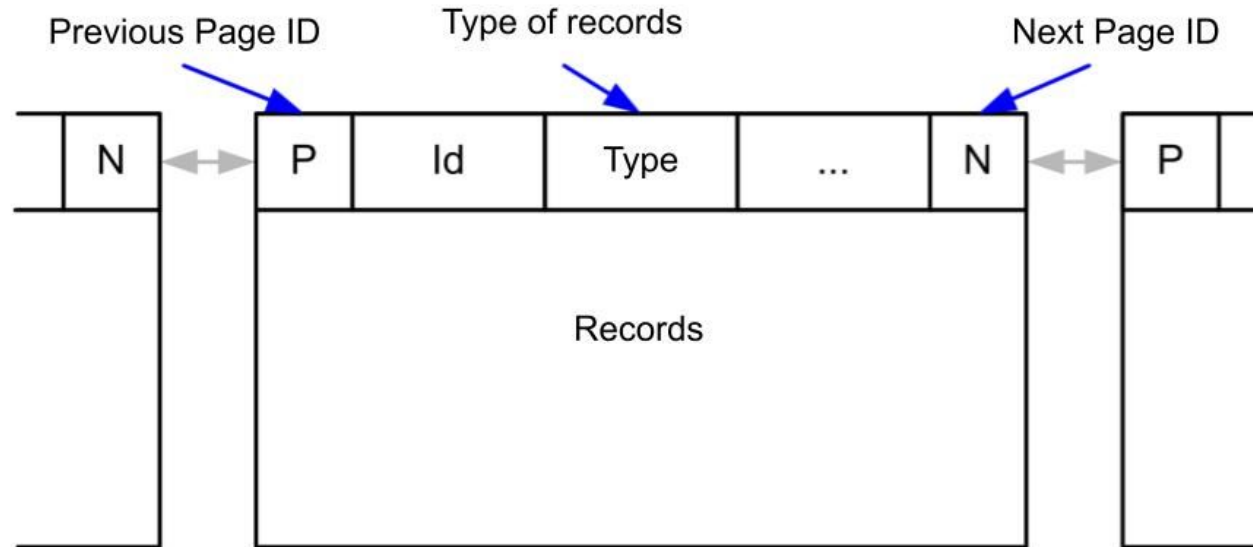
CITY UNIVERSITY  
LONDON

- Data is stored in files
- A file can contain several tables
- A table consists of several pages
- Page contains several records

# Records management



CITY UNIVERSITY  
LONDON



# Record ID

---



CITY UNIVERSITY  
LONDON

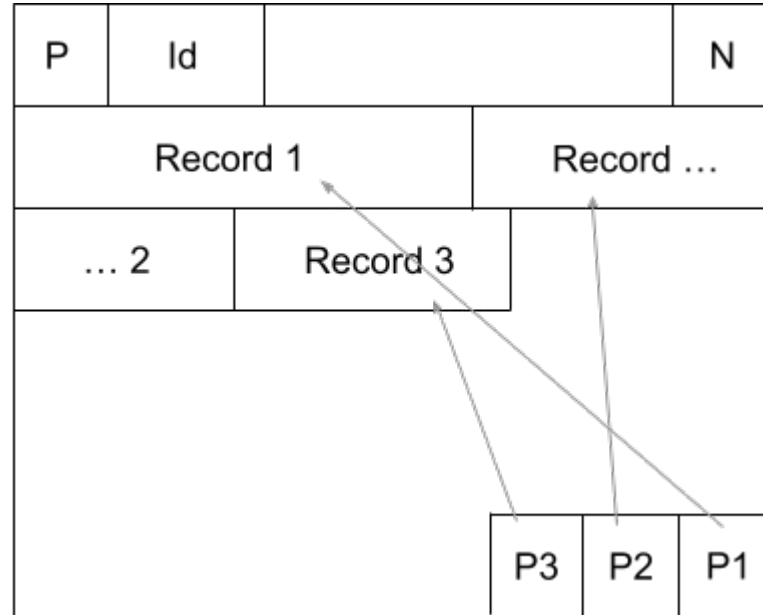
- Record ID (RID) consists of two parts:
  - Page ID
  - ID Record on the page
- We do not want it to change through time





# Records on page



CITY UNIVERSITY  
LONDON



# Records Overflow



# Indexing

---



CITY UNIVERSITY  
LONDON

- There can be millions of records
- How to find the one necessary?
  - Search through the whole table
  - Have an additional data structure Index

# Index

---




CITY UNIVERSITY  
LONDON

- Some subset of columns
- Not necessarily primary key
- Build at first and then updated
- Fast search gives the pointer to the record

# Simple index



CITY UNIVERSITY  
LONDON



Why we do not want to store already in the sorted way: (Key, Record)?

# Hash Index



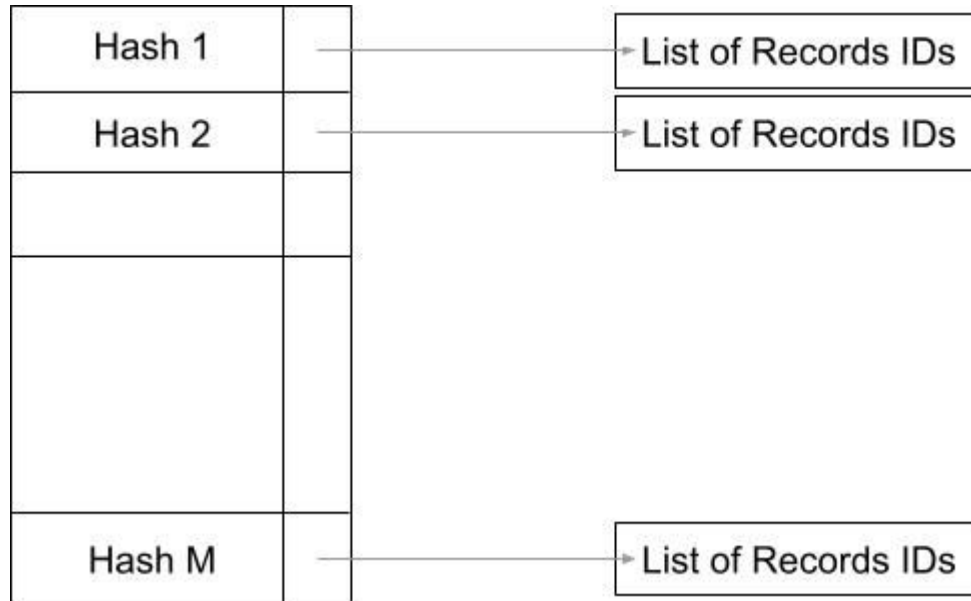
CITY UNIVERSITY  
LONDON

- Each defined index tuple is transformed into some integer
- There is a table that given some integer returns a list of elements
- HashMap translated a tuple into integer and look into the corresponding element in the table
- Expected number of elements with the same hash and different values is  $O(1)$

# Hash Table



CITY UNIVERSITY  
LONDON



# Improved Queries

---



CITY UNIVERSITY  
LONDON

- Key uniqueness
- IN
- EXISTS
- COUNT
- JOIN





# Ordered Index

---

- The keys are sorted
- Search is like searching in the sorted sequence
- Can support range queries

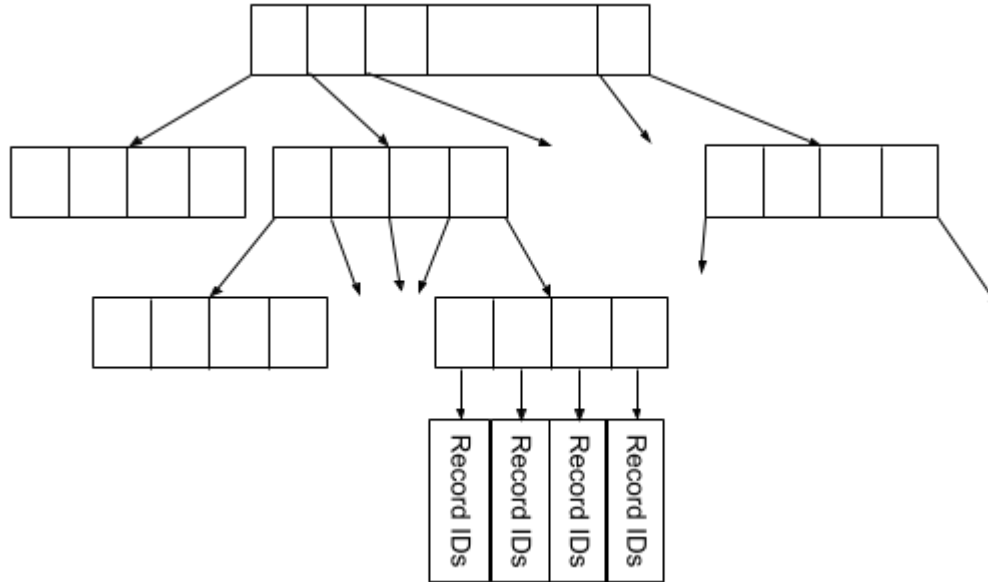
For that we need a search tree.

In databases, typically, B-tree is used

# BTree



CITY UNIVERSITY  
LONDON





# Multidimensional Index



CITY UNIVERSITY  
LONDON

- Suppose that the key of the index consists of several attributes
- One can use R-index
- It splits the dimension into cuboids



- Usually, already exists for primary keys (how to check uniqueness?)
- Need to choose the representative ones with good selection
  - Surname is fine
  - Age is maybe ok
  - Gender is not really helpful
- The number of elements per each “bucket” should be not large
- Database can use some statistics to not use the asked index

# Suggestions

---



CITY UNIVERSITY  
LONDON

- On keys
- On foreign keys
- For range queries
- For join operations

# MySQL syntax

---



CITY UNIVERSITY  
LONDON

```
CREATE INDEX index_name {BTREE | HASH}  
ON tbl_name (key_part,...);
```

# Sources

---



CITY UNIVERSITY  
LONDON

C.J. Date, An Introduction to Database Systems

D. Knuth, The Art of Computer Programming, Volume 3, Sorting and Searching

A. Silberschatz, H. F. Korth, S. Sudarshan, Database System Concepts