

Concurrency Control

Concurrency Control



CITY UNIVERSITY
LONDON

- Concurrency control: ensure that one user's work does not inappropriately influence another user's work
- No single concurrency control technique is ideal for all circumstances
- Trade-offs need to be made between level of protection and throughput

Atomic Transactions



CITY UNIVERSITY
LONDON

- A transaction, or logical unit of work (LUW), is a series of actions taken against the database that occurs as an atomic unit
- Either all actions in a transaction occur - **COMMIT**
- Or none of them do - **ABORT**

Example



CITY UNIVERSITY
LONDON

The common example is the transfer of money

```
read(ACC1);  
ACC1 := ACC1 - 50;  
write(ACC1); read(ACC2);  
ACC2 := ACC2 + 50;  
write(ACC2);
```

Transaction Outcomes & Boundaries



CITY UNIVERSITY
LONDON

- Database transaction - a set of operations of these kinds:
 - Transaction boundary/edge: begin, commit, abort (rollback)
 - Reads (SELECT), writes (DELETE, INSERT, UPDATE)
- Committed: when it completes successfully;
- Aborted (Rolled-back): when it is not executed successfully;

Transaction properties:

- Atomicity: either all the operations in a transaction are executed against the database, or none is.
- Consistency: the execution of a transaction must preserve the consistency of the database
- Isolation: the execution of a transaction must preserve the consistency of the database.
- Durability: the effects of a successfully completed transaction are stored in the database permanently (not lost due to HW/SW failure)

Concurrent Transaction



CITY UNIVERSITY
LONDON

- Concurrent transactions: transactions that are being processed at the same time
- Can be also run on one CPU
- Transactions are interleaved
- Concurrency problems
 - Lost updates
 - Inconsistent reads

Concurrency Issues



CITY UNIVERSITY
LONDON

Transfer(A, B, 100)

1. mA = read(moneyA) // 1000

2. mB = read(moneyB) // 1000

6. write(moneyA, mA - 100)

7. write(moneyB, mB + 100) // 1100

Transfer(B, A, 100)

3. mB = read(moneyB) // 1000

4. mA = read(moneyA) // 1000

5. write(moneyB, mB - 100)

8. write(moneyA, mA + 100) // 1100

Concurrency Control



CITY UNIVERSITY
LONDON

- A definition: The mechanism of managing simultaneous operations (executed by multiple transactions) on the database preventing them to interfere with one another.
- How can the system control the interaction among, and effects of, concurrent transactions?
- Different CC mechanisms exist (in different DBMSs)
- If control of execution of concurrent operations is left to the underlying OS, many possible schedules, i.e. orders of operation executions, are possible
- Including the ones that would leave DB in an inconsistent state, and/or the ones that would send wrong results to DB users.



Transaction isolation levels

- The Isolation levels control the degree to which the execution of one transaction is isolated from the execution of other concurrent transactions (accessing the same data).
- A lower isolation level increases concurrency (improves performance), but can impair data correctness. Conversely, a higher isolation level ensures that data remains correct, but can negatively affect concurrency.
- ISO SQL-92 defines the following isolation levels (from the strictest to the most relaxed):
 - Serializable, Repeatable read, Read committed and Read uncommitted
 - Based on the defined anomalies: Dirty Read, Non-Repeatable and Phantom Reads.

Isolation Levels



CITY UNIVERSITY
LONDON

Isolation Level	Transaction s	Dirty Reads	Non-Repeatable Reads	Phantom Reads
TRANSACTION_NONE	Not supported	Not applicable	Not applicable	Not applicable
READ_UNCOMMITTED	Supported	Allowed	Allowed	Allowed
READ_COMMITTED	Supported	Prevented	Allowed	Allowed
REPEATABLE_READ	Supported	Prevented	Prevented	Allowed
SERIALIZABLE	Supported	Prevented	Prevented	Prevented

Serializability



CITY UNIVERSITY
LONDON

- Transaction Schedule: Sequence of operations by concurrent transactions that preserves the order of the operations in each individual transaction
- Serial Schedule: The transactions are performed in serial order (operations are not interleaved) sequentially with no overlap in time
- A transaction schedule is serializable if there exists a serial schedule leading to the same state of the database
- A database transaction manager is serializable if this happens for all schedules

Locks



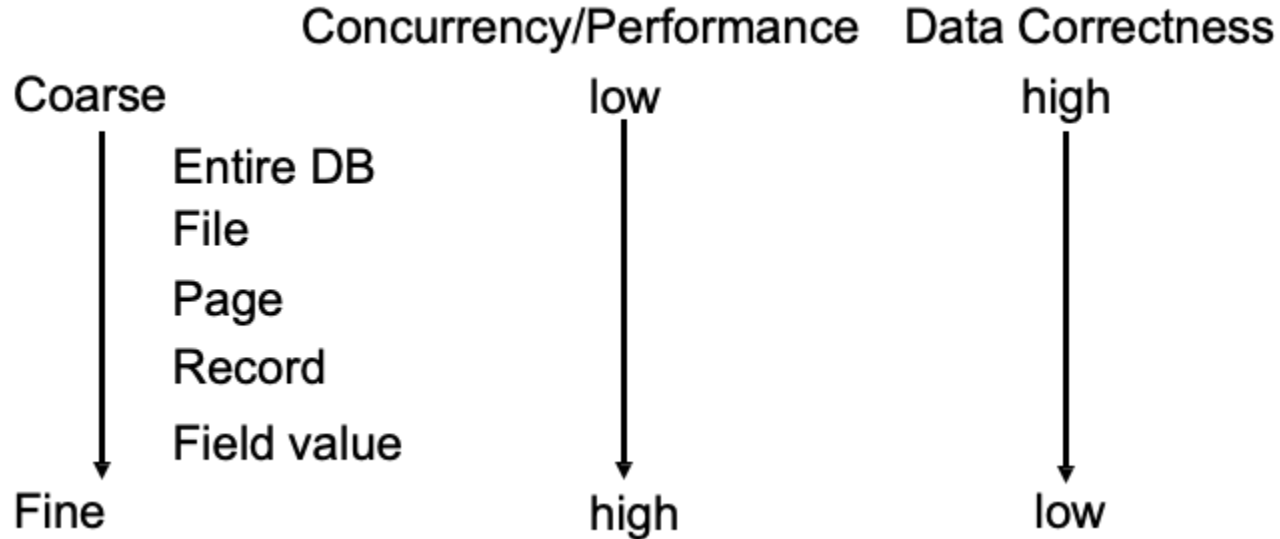
CITY UNIVERSITY
LONDON

- How can we guarantee serializability?
- We can use locks on the data
- It controls the access to the data
- Two lock types:
 - Exclusive (write) lock
 - Shared (read) lock

Lock Granularity trade-offs



CITY UNIVERSITY
LONDON



Two-Phase Locking Protocol (2PL)



CITY UNIVERSITY
LONDON

- Phase 1 (Growing/expanding phase)
 - A transaction T always acquires a lock on an object before reading (S-lock), or writing (X-lock), it.
- Phase 2 (Shrinking phase)
 - After releasing a lock, transaction T cannot acquire a new lock.

There is impossible to get lock after the first unlock.

Several variations: Conservative 2PL, Strict 2PL, Strong Strict 2PL

Original 2PL takes a lock when necessary

Conservative 2PL takes all at the beginning

Deadlock



CITY UNIVERSITY
LONDON

When transactions needs some resource that is already locked and it is impossible to release anything.

Transfer(A, B, 100)

1. lock(accountA)
2. accountA += 100
5. lock(accountB) // cannot

Transfer(B, A, 100)

3. lock(accountB)
4. accountB += 100
6. lock(accountA) // cannot

Deadlock. Consequences.



CITY UNIVERSITY
LONDON

- It cannot happen in Conservative PL. It just takes the locks in some proper order.
- To find it, the database uses dynamic ownership graph.
- To solve it, one of the transactions is rolled back and timeouted.
- There is some logic which decides which one to roll back.

Sources



CITY UNIVERSITY
LONDON

C.J. Date, An Introduction to Database Systems (Sections 15 and 16)

J. Ullman and J. Widom, A First Course in Database Systems (Section 6.6)

Bernstein P.A., Hadzilacos V., Goodman N. Concurrency Control and Recovery in Database Systems