

# IN2029: Programming in C++

## Session 2 – Sequential containers

Vahid Rafe

Department of Computer Science  
City, University of London

Autumn term, 2025



# This session

We'll be writing some programs that operate on batches of data, which allows us to explore

- a bit more about streams
- the standard idiom for looping to the end of an input stream
- manipulators
- vectors from the standard template library
- introduction to containers

# Calculating statistics from a list of numbers

**Task:** read in a list of numbers and print their average.

The overall structure of our program will be:

```
#include <iostream>
#include <iomanip>

using namespace std;

int main() {
    // ... read in data ...
    // ... print results ...
    return 0;
}
```

## Reading the data

The first part is to read all the numbers and record their count and sum:

```
cout << "Please enter a series of numbers\n";

// the number and total of values read
int count = 0;
double sum = 0;

// read values from standard input
double x; // a variable for reading into
while (cin >> x) {
    ++count;
    sum += x;
}
```

## Library details: testing for end-of-input

We have already seen that the `>>` operator returns the input stream, in statements like

```
cin >> x >> y >> z;
```

The result of `>>` can also be used in a test, as in the common idiom for reading a series of things and testing for the end of the input:

```
while (cin >> x) {  
    // .. do something with x  
}
```

Testing a stream yields `true` if the last operation on the stream succeeded, and `false` if it didn't.

This idiom is much more robust than separately testing for eof.

(You can indicate end of input on the console by typing Control-Z Return on Windows, or Control-D on Unix.)

## Alternative: reading from a file

If we want to read from a file instead of the standard input, we use an `ifstream` (a different kind of `istream`). At the top of the file, we need to include another header:

```
#include <fstream>
```

Then we declare a variable `in` of type `ifstream`, with the file name as a parameter, and then read from that instead of `cin`:

```
// read values from a file
ifstream in("values.txt");
double x;
while (in >> x) {
    ++count;
    sum += x;
}
```

## Breaking the input into words

An example reading strings:

```
#include <string>
#include <iostream>

using namespace std;

int main() {
    string s;
    while (cin >> s)
        cout << s << '\n';
    return 0;
}
```

Recall that the >> operator on strings reads words.

## reading a line including spaces

If we want to read a string containing spaces from standard input or from a file, we use a `getline` function. At the top of the file, we need to include this header:

```
#include <iostream>
```

For reading from a file, use `getline(in, line)`, for reading from standard input use `getline(cin, line)`, you need to define `line` as a `string` and `in` should be defined as an `ifstream`:

```
// read values from a file
ifstream in("values.txt");
string line;
while (getline(in, line)) { //or you can use getline(cin, line)
    //...
}
```



## Language details: `i++` vs `++i`

The following statements all increase an `int` variable `i` by one:

```
i = i+1;  
i += 1;  
i++;  
++i;
```

The difference between the last two is only seen when the value of the expression is used:

```
int i = 5;  
int j = ++i; // j is set to 6; i is now 6  
int k = i++; // k is also set to 6; i is now 7
```

- `i++` returns the value **before** incrementing (so the old value has to be saved somewhere, which could be expensive with some types)
- `++i` returns the value **after** incrementing (simpler)

## Printing the results

Finally, we want to print the results:

```
cout << count << " numbers\n";  
if (count > 0) {  
    cout << "average = " << sum/count << '\n';  
}
```

By default, floating point numbers are printed with up to 5 significant figures, but we can change that:

```
cout << "average = " << setprecision(3) <<  
    sum/count << '\n';
```

## Library details: manipulators

`setprecision(3)` is an example of a stream **manipulator** (from the `<iomanip>` system header), like `flush` or `endl`: a special kind of object with an overloading of the `<<` operator that changes the state of the stream.

This manipulator is used to adjust formatting:

```
cout << setprecision(3);
```

doesn't do any output, but it sets the precision for any following output.

```
cout << setprecision(3) << x << setprecision(5) << y;
```

Other manipulators set base, paddings, etc.

## Cleaning up

- We have used `setprecision` to set the maximum number of significant figures to what we want.
- Nothing else is happening in this program, but in general it would be polite to set the precision back to what it was before.
- We can get the current precision using `cout.precision()`.

This yields our final version:

```
int prec = cout.precision();  
cout << "average = " << setprecision(3) <<  
    sum/count << setprecision(prec) << '\n';
```

## Calculating a different statistic

**Task:** read in a list of numbers and print their median.

The **median** of a collection of numbers is the “middle” value when they are arranged in order:

1 3 3 7 10 11 11 13 14 15 15

However, the input data may be in any order.

- Unlike computing the average, to compute the median we will need to store all the numbers until the end of the program. We shall use a **vector** to do this.
- Then we need to arrange the values in order. We shall use the library function **sort**.
- Then the median will be the middle value in the vector.

# Outline

The overall structure of our program will be:

```
#include <iostream>
#include <vector>
#include <algorithm>

using namespace std;

int main() {
    // ... read and store the data ...
    // ... sort the data ...
    // ... print the middle value ...
    return 0;
}
```

# Vectors

```
#include <vector>
```

C++ has arrays, but we'll use vectors instead (a container like `ArrayList` in Java, except that a variable of `vector` type holds an object, not a reference):

```
vector<int> vi;    // empty vector of ints  
vector<string> vs; // empty vector of strings
```

Vectors also be extended:

```
vs.push_back(s);
```

The current length of `vs` is `vs.size()`

Vectors can be accessed just like arrays (indices `0 ... size() - 1`):

```
vi[1] = x;  
vi[2] = vi[1] + 3;
```

# Implementation of vectors

Internally, vectors are implemented as extensible arrays (see session 3 of Data Structures and Algorithms). One possible implementation is:



The precise internals do not concern us here.

We need only use the provided operations, confident that they are *very* carefully engineered.



## Reading the data into a vector

We start by reading all the numbers and storing them in a vector:

```
cout << "Please enter a series of numbers\n";

// read numbers from the standard input
// and store them in a vector
vector<double> v;
double x;
while (cin >> x)
    v.push_back(x);
```

We don't need a separate variable to count them: we can use `v.size()`.

## Finding the median: outline

- Only a non-empty vector can have a median.
- First, we need to sort the vector.

```
// compute and output results
unsigned n = v.size();
cout << n << " numbers\n";
if (n > 0) {
    // sort the whole vector
    sort(v.begin(), v.end());

    // ... find the middle value
}
```

## Language details: unsigned types

C++ has signed and unsigned integral types of various sizes:

Signed	?	Unsigned
<code>signed char</code>	<code>char</code>	<code>unsigned char</code>
<code>short</code>		<code>unsigned short</code>
<code>int</code>		<code>unsigned int</code> (or <code>unsigned</code> )
<code>long</code>		<code>unsigned long</code>
<code>long long</code>		<code>unsigned long long</code> (in C++11)

- Unlike in Java, the sizes are not defined by the standard (but they are non-decreasing).
- `char` may be either a signed or unsigned type, whichever is more efficient on this architecture.
- Unsigned types cannot be negative: if `i` is of unsigned type, `i < 0` can never be `true`.

## Unsigned types: caution

- Unsigned integers will silently underflow:

```
unsigned i = 0;  
i--;
```

will not fail – it will set `i` to a very large positive number.

- If an operation involves both a signed and unsigned type, it will silently convert the signed type to unsigned first, so in

```
int i = -5;  
unsigned j = 1;  
if (i < j)
```

the last test will fail, because `-5` will be silently converted to a very large positive number.

## The type of `size()`

- Containers cannot have negative size.
- The return type of the `size()` member function is an unsigned type, but *which* unsigned type is implementation dependent.
- The portable name of its type is `vector<double>::size_type`.
- Here `::` selects a static attribute of the type `vector<double>`. (This is a different use of `::` from namespace qualification, as in `std::vector`.)
- We can use this as the type of the variable `n`:

```
vector<double>::size_type n = v.size();
```

## Library details: **sort**, **begin**, **end**

```
sort(v.begin(), v.end());
```

- To sort a vector, we use the **sort** function, declared in the `<algorithm>` system header.
- Instead of a container, **sort** takes two positions or **iterators** (which we'll explore in session 4).
- These positions should be in the same container, with the first before the second.
- The vector class has member functions **begin()** and **end()**, yielding positions as the start and end of the vector.
- So the above statement sorts the whole vector – a common idiom, but using iterators is more general.

# Where is the median?

There are two cases:

- odd number of elements, e.g. 9:

	0	1	2	3	4	5	6	7	8
<b>v</b>									

middle element is cell 4, i.e. `v[v.size()/2]`

- even number of elements, e.g. 8:

	0	1	2	3	4	5	6	7
<b>v</b>								

In this case we average the two middle elements (cells 3 and 4):

$$(v[v.size()/2 - 1] + v[v.size()/2]) / 2$$

# Computing the median

We use this plan to compute the median of the sorted array:

```
// find the middle value
vector<double>::size_type middle = n/2;
double median;
if (n%2 == 1) // size is odd
    median = v[middle];
else // size is even
    median = (v[middle-1] + v[middle])/2;
cout << "median = " << median << '\n';
```

and our program is complete.



## Type aliases

In C++11, we can declare a new name for a type: A **typedef** declaration allows us to introduce a new name for a type:

```
using vec_size = vector<double>::size_type;
```

This defines a new type name **vec\_size** that is equivalent to the longer name. One use is to avoid repeating a long type name:

```
vec_size n = v.size();  
// ...  
vec_size middle = n/2;
```

Older versions of C++ used the **typedef** declaration, with exactly the same effect:

```
typedef vector<double>::size_type vec_size;
```

## C++11 shortcut: `auto`

In C++, we can just use `auto` instead of the type, which tells the compiler that the variable has the type of the initializing expression:

```
auto n = v.size();  
// ...  
auto middle = n/2;
```

Here `n` has type `vector<double>::size_type`, and `middle` also has an unsigned type.

- A variable declared as `auto` must be initialized.
- The variable still has a type, which determines how it can be used, and what operations on the variable mean, it's just implicit.

Prefer `auto` to explicit type declarations (if there is an initializing expression with the desired type).

## Vectors: further points

- A vector variable contains a whole vector:

```
vector<int> v1 = v; // copy the vector  
sort(v.begin(), v.end());
```

results in `v` being sorted, but `v1` still containing a copy of the original unsorted `v`.

- When indexing `v[i]`, the index `i` is not checked: if it is out of range, the program may crash or continue with corrupted data.
- Other vector member functions:
  - `back()` returns the last element of the vector
  - `pop_back()` removes the last element of the vector

## Another container: **deque**

Deque (double-ended queues) can be created in a similar way:

```
deque<int> d; // an empty deque
```

Deque support indexing with `[]`, and these member functions:

`size()` the number of elements in the deque

`push_back(x)` add `x` to the back of the deque

`back()` returns the last element of the deque

`pop_back()` removes the last element of the deque

`push_front(x)` add `x` to the front of the deque


`front()` returns the first element of the deque

`pop_front()` removes the first element of the deque

There are common names with **vector**, but no inheritance.

# Implementation of deques

Deques are also implemented as extensible arrays, but unlike vectors their front end can also move. One possible implementation is:



Once again, we just use the provided operations, relying on the carefully tuned implementation.

## Another container: **stack**

Stacks can be created in a similar way:

```
stack<int> s; // an empty stack
```

Stacks do not support indexing with `[]`, but support these member functions:

`size()` the number of elements in the stack

`push(x)` add `x` to the top of the stack

`top()` returns the last element (on top) of the stack

`pop()` removes the last element of the stack

Queues are very similar to stacks: (`size()`, `push(x)`, `front()` and `pop()`)

```
queue<int> q; // an empty queue
```

## Next week

- Functions in C++ allow us to structure and reuse code.
- Passing parameters by value (like in Java) involves copying, which can be expensive as in C++ (unlike in Java) variables contain whole objects.
- Passing parameters by reference avoids copying, and is heavily used in C++.
- It is good practice to use `const` qualifiers to declare that you're not changing something.