

IN2029: Programming in C++

Session 3 – Program organization

Vahid Rafe

Department of Computer Science
City, University of London

Autumn term, 2025



Refactoring

Refactoring is re-arranging code (to improve re-use, maintainability, *etc*), without changing what it does. We shall use:

- Functions, with
 - parameters passed by value
 - parameters passed by `const` reference
 - parameters passed by non-`const` reference
- Using multiple source files.
 - source files and header files
 - dealing with repeated inclusion

Functions

Consider the program that computes medians and averages.

- Several pieces could be made into re-usable components:
 - read a vector
 - compute the median of a vector
 - compute the average of a vector
- We will split these out as separate functions.
- In C++, functions must be declared before use, so our new functions will go after the `using` declaration, but before `main()`.
- The median and average functions won't do any I/O: this makes them more generally useful.

Median of a vector

```
// the median of the values in a vector
// requires: v.size() > 0
double median(vector<double> v) {
    auto n = v.size();
    // sort the whole vector
    sort(v.begin(), v.end());
    auto middle = n/2;
    if (n%2 == 1) // size is odd
        return v[middle];
    else // size is even
        return (v[middle-1] + v[middle])/2;
}
```

Checking the precondition

- The `median` function promises to return the median of a vector, but only if the vector is non-empty.
- This is the **precondition** on the function.
- If the precondition is not satisfied, the function can do anything. Here we throw an exception `std::domain_error` (from the `<stdexcept>` header).

```
// the median of the values in a vector
// requires: v.size() > 0
double median(vector<double> v) {
    auto n = v.size();
    if (n == 0)
        throw domain_error("median of an empty vector");
    // ...
}
```

Parameter passed by value

```
double median(vector<double> v) {  
    sort(v.begin(), v.end()); // rearranges v  
    // ...  
}  
  
int main() {  
    vector<double> values;  
    // ...  
    cout << "median = " << median(values) << '\n';  
    // ... (values unchanged)  
}
```

- The parameter `v` is a new vector, initialized as a **copy** of `values`.
- Changes to `v` (e.g. `sort`) do not affect `values`.

Avoiding the copy

- In **median**, it makes sense to operate on a copy of the vector, because the function changes the vector (by sorting it), and the rest of the program might require the values in the original order.
- But copying large arguments (e.g. vectors) is expensive, and often we want to avoid it.
- Also, we might want to return data through the argument, and we can't do that if we're operating on a copy.
- A solution (Fortran, Pascal, C++, C#, etc) is **reference parameters**.
- In C++, reference parameters are marked with **&**.
- If the function is not going to change the argument, it can be marked as **const**.

Prefer pass-by-reference-to-const to pass-by-value (except for primitive types).

Average of a vector

When computing the average, we can use the original vector (no copying), and not change it:

```
// the average of the values in a vector
// requires: v.size() > 0
double average(const vector<double> &v) {
    auto n = v.size();
    if (n == 0)
        throw domain_error("average of an empty vector");
    double sum = 0;
    using vec_size = vector<double>::size_type;
    for (vec_size i = 0; i < n; ++i)
        sum += v[i];
    return sum / v.size();
}
```


Parameter passed by const reference

```
double average(const vector<double> &v) {  
    // ... (doesn't change v)  
}  
  
int main() {  
    vector<double> values;  
    // ...  
    cout << "average = " << average(values) << '\n';  
    // ... (values unchanged)  
}
```

- The parameter **v** is another name for **values** (no copying).
- **const** promises that the function will not change **v**, and the compiler checks this.

Other uses of `const`

The `const` specifier can also be used

- to declare global constants, e.g.

```
const int days_per_week = 7;
```

(This is much better than using C-style `#defines`.)

- to declare parameters and local variables with fixed values, e.g.

```
const auto middle = v.size() / 2;
```

- on member functions in classes (session 6)

Use `const` whenever possible.

Another example of `const`

For efficiency, we might save `v.size()` in a variable `n`, which we then do not change:

```
// the average of the values in a vector
// requires: v.size() > 0
double average(const vector<double> &v) {
    const auto n = v.size();
    if (n == 0)
        throw domain_error("average of an empty vector");
    double sum = 0;
    using vec_size = vector<double>::size_type;
    for (vec_size i = 0; i < n; ++i)
        sum += v[i];
    return sum / n;
}
```

A function that modifies a parameter

Because a reference parameter is an alias for the argument, changing it modifies the argument directly, e.g.

```
// add one to each element of a vector
void add_one(vector<double> &v) {
    using vec_size = vector<double>::size_type;
    for (vec_size i = 0; i < v.size(); ++i)
        v[i]++;
}

int main() {
    vector<double> values;
    // ...
    add_one(values); // changes values
}
```

Other uses of references

We can also use references when we want to re-use the same memory location (rather than re-using a value), e.g.:

```
vector<double> v;  
// ...  
double &x = v[12];  
x = 1 + 2*x;
```

Here **x** is not a new variable, but an alias for **v[12]**. Everything that is done to **x** is actually done to **v[12]**.

- C++ references have no counterpart in Java.
- Java references are similar to C++ pointers (session 9).


Function to read a vector of numbers

This function also operates on the original input stream (no copying), but changes it so the parameter `in` is passed by reference, but without `const`:

```
// read numbers from an input stream
// and return them in a vector
vector<double> read_vector(istream &in) {
    vector<double> v;
    double x;
    while (in >> x)
        v.push_back(x);
    return v;
}
```

Why is returning a vector cheap?

Recall the implementation of a vector (details may vary):



- Normally if you assign a vector, both parts get copied.
- But here the vector is in a local variable, which is about to disappear, so the compiler knows that only the record needs to be copied, not the array. (The vector is *moved* rather than copied.)
- A particularly smart compiler might even avoid copying the record.

Main program

```
int main() {  
    cout << "Please enter a series of numbers\n";  
  
    // read numbers from the standard input  
    // and store them in a vector  
    const auto values = read_vector(cin);  
  
    // compute and output results  
    const auto n = values.size();  
    cout << n << " numbers\n";  
    if (n > 0) {  
        cout << "average = " << average(values) << '\n';  
        cout << "median = " << median(values) << '\n';  
    }  
    return 0;  
}
```


Three kinds of parameters

- parameter passed by value:

```
double median(vector<double> v)
```

The parameter is a **new variable**, initialized as a copy of the actual parameter.

- parameter passed by **const** reference:

```
double average(const vector<double> &v)
```

The parameter is an **alias** for the actual parameter, but the function will not change it.

- parameters passed by non-**const** reference:

```
vector<double> read_vector(istream &in)
```

The parameter is an **alias** for the actual parameter, and the function may change it.

Return by reference

- Normally a function returns its result by value (so it is copied).
- It is possible to define a function that returns by reference, e.g. the `<iostream>` library function to read a line:

```
istream& getline(istream& in, string &s)
```

This returns its parameter `in`.

- Then the function call is an alias for whatever is returned, e.g. in

```
string s;  
while (getline(cin, s))  
    cout << s.size() << '\t' << s << '\n';
```

- You can't return a temporary or local variable by reference (because they will be gone when the function returns).

Separate compilation


- For large projects, it is common to split programs into several files, which can be compiled separately as required.
- In addition to the source files (`.cpp` or `.cc`, *etc*), we put declarations in **header files** (`.h`).
- Typically, `foo.h` contains declarations for things defined in `foo.cpp`.
- source files (and other header files) access these definitions using **`#include`** lines, which insert the text of those header files at compile time:

```
#include "foo.h"
```

Note the different syntax from system headers.

- In particular, `foo.cpp` should include `foo.h` (so the compiler checks that the definitions match the declarations).

Typical program structure



Rebuilding a program after edits

- If a source file has been changed, it must be recompiled, generating the corresponding object file.
- If a header file has been changed, all the source files that include it, directly or indirectly, must be recompiled.
- If any source files have been recompiled, the object files must be linked to regenerate the executable program.

Minimize compilation dependencies between files.

- Minimize the number of header files included in other files.
- This also means putting the minimum in each header file.

Splitting our program

We shall re-arrange our program to extract an embryonic statistics library. The new program will consist of the following files:

- `stats.h` A header file, with declarations of the functions `median` and `average` (just their parameters and return types)
- `stats.cpp` A source file that starts with `#include "stats.h"` and gives the full definitions of the functions `median` and `average`.
- `main.cpp` A source file that starts with `#include "stats.h"` and gives the full definitions of the functions `read_vector` and `main`.

The auxiliary source file `stats.cpp`

```
#include "stats.h"
#include <vector>
#include <algorithm>
#include <stdexcept>

using namespace std;

// the median of the values in a vector
// requires: v.size() > 0
double median(vector<double> v) {
    const auto n = v.size();
```

and so on (rest of `median` and all of `average`).

Placement of `#includes`

```
#include "stats.h"  
#include <vector>  
#include <algorithm>  
#include <stdexcept>  
  
using namespace std;
```

- We include our own header files before system headers, which come before any `using` declaration. That ensures that our header files don't assume anything had been included.
- We only include the system headers we need. (In this case, we don't need `<iostream>`.)

The header file `stats.h`

```
#include <vector>

// the median of the values in a vector
// requires: v.size() > 0
double median(std::vector<double> v) ;

// the average of the values in a vector
// requires: v.size() > 0
double average(const std::vector<double> &v) ;
```

- Never put `using` declarations in header files – that would force names on clients.
- This means all standard names must be fully qualified.
- Give only function signatures, followed by ‘;’, not the bodies.
- For the types, we only need the `<vector>` system header.

The main source file `main.cpp`

- The file begins with the headers, following the same principles as in `stats.cpp`:

```
#include "stats.h"  
#include <iostream>  
#include <vector>  
  
using namespace std;
```

- This is followed by the functions `read_vector` and `main`.
- The `main` function uses `median` and `average`, which are declared in `stats.h`.

Handling repeated inclusion

- Sometimes a header file can be included twice, e.g.:

```
#include "line.h"      // includes point.h
#include "polygon.h"   // also includes point.h
```

- With just function declarations this is harmless, but with later features it will cause compilation errors.
- We need to ensure that the second inclusion of any header file does nothing.
- A widely-implemented, but non-standard method is to place this line at the top of the header file:

```
#pragma once
```

- The fully portable method uses **include guards**.

stats.h with include guards

```
#ifndef GUARD_stats_h
#define GUARD_stats_h

#include <vector>

// the median of the values in a vector
// requires: v.size() > 0
double median(std::vector<double> v) ;

// the average of the values in a vector
// requires: v.size() > 0
double average(const std::vector<double> &v) ;

#endif
```

Include guards

```
#ifndef GUARD_stats_h
#define GUARD_stats_h

// body of header file

#endif
```

- The `#` lines are features from the C preprocessor that shouldn't be used for anything else.
- The first time the file is included, `GUARD_stats_h` is not defined, so the whole is included, including the second line, which defines `GUARD_stats_h`.
- The next time the file is included, `GUARD_stats_h` is defined, so the rest of the file is skipped.