

Session 4

Iterators and associative containers

Containers

- A *container* (called a “collection” in Java) is an object holding a collection of values.
 - A *sequential container* (e.g. **vector**, **deque**, **list**) holds a collection of values arranged in order.
 - An *associative container* (e.g. **map**) holds values that can be efficiently looked up by key.
- An *iterator* is an object representing a position within a container.
- Functions in the standard **<algorithm>** library tend to operate on iterators rather than containers (more next week).

Outline

We will introduce iterators by going through three versions of a program, all doing the same thing:

1. operating on a vector, using indices (numbers) to refer to positions in the vector
 - correct but inefficient, due to the limitations of vectors.
2. still operating on a vector, but using iterators to refer to positions in the vector
 - equivalent to version 1, but now easier to switch to a different container.
3. operating on a linked list, but iterators to refer to positions
 - same code as version 2, but now more efficient because it is operating on a linked list.

Using indices

Printing out a vector

One of the exercises from last week was to write a function to print a vector. Here is a solution:

```
// write vector to an output stream
void write_vector(ostream &out, const vector<double> &v) {
    const auto n = v.size();
    out << "vector:";
```

```

using vec_size = vector<double>::size_type;
for (vec_size i = 0; i < n; ++i)
    out << ' ' << v[i];
out << '\n';
}

```

Deleting all the zeroes from a vector

Task: write a function that deletes all the zeroes from a vector.

- The signature of our function will be

```
void delete_zeroes(vector<double> &v)
```

- To delete elements, we shall use the member function **erase**:

```
v.erase(v.cbegin() + i);
```

- The argument of **erase** is an *iterator*, an object representing a position in the vector container.
- Here **v.cbegin()** represents the position at the start of the vector, and **v.cbegin() + i** represents the position **i** places further on.

The deletion function

```

// remove zero values from a vector
void delete_zeroes(vector<double> &v) {
    using vec_size = vector<double>::size_type;
    vec_size i = 0;
    while (i < v.size())
        if (v[i] == 0)
            v.erase(v.cbegin() + i);
        else
            ++i;
}


```

⚠ Caution

We cannot optimize this by putting **v.size()** in a variable, because **erase** changes the size of the vector.

A performance problem

- erase** can be expensive for vectors if the index is not at the end of the vector.
- Action of **erase** at position **i**:



- All elements after the erased position must be moved down.
- As a result, `delete_zeroes` can take $O(n^2)$ time.
- We could fix this particular function with a different algorithm, but instead we'll switch to a different data structure (a linked list) for which `erase` is efficient.
- But first, we generalize from indices (which work with vectors) to iterators (which work with any container).

Iterators

Random versus sequential access

To make our code work with a different data structure, we need to think about what primitives we really need. The fewer we require, the more options we will have.

- Vectors provide a powerful indexing operation `v[i]` (*random access*).
- For the two functions we want to write, it is sufficient to step through the elements one by one from the start of the container. (*sequential access*).
- The operations we need are:
 - access the current element
 - move to the next element
 - test whether we've reached the end of the container
- There are many structures (*e.g.* linked lists) that will support these operations but not indexing.

Iterators

- An iterator is an object representing a position in a container (*e.g.* a vector).
- Iterators for different container types are typically implemented differently.
- For any iterator `it`, the following operations are supported:
 - `*it` is a reference to the element currently pointed to by the iterator.
 - `++it` moves the iterator on to the next position.
 - `it == it2` or `it != it2`, where `it2` is another iterator for the same container, tests whether the two iterators are at the same position or not.
- Iterators on some container types support more operations (*e.g.* the `+` operation on vector iterators that we've already seen).

Iterator types for a container

Like every standard container, `vector<double>` defines two iterator types:

- `vector<double>::iterator`
- `vector<double>::const_iterator`

The difference is that for `const_iterator`, `*it` is a `const` reference, so that

```
*it = 1.2;
```

is only allowed for `iterator`.

☞ *Advice (Effective Modern C++ Item 13)*

Prefer `const_iterator`s to `iterator`s (where possible).

(for the same reasons we use `const` wherever possible)

Standard iterators for a container


Every standard container provides the following member functions:

`begin()` returns an `iterator` pointing at the first element of the container.

`end()` returns an `iterator` pointing just after the last element of the container.

`cbegin()` returns a `const_iterator` pointing at the first element of the container. (C++14)

`cend()` returns a `const_iterator` pointing just after the last element of the container. (C++14)

**Looping through a container**

You can't use `*` on `end()` or `cend()`, but you can use `==`.

The standard idiom, when not changing the elements:

```
for (auto it = v.cbegin(); it != v.cend(); ++it)
    // do something with *it
```

(Here the type of `it` is `vector<double>::const_iterator`.)

When changing the elements:

```
for (auto it = v.begin(); it != v.end(); ++it)
    *it += 3;
```

(Here the type of `it` is `vector<double>::iterator`.)

Note on language versions

The above is for C++14. Earlier versions of C++ did not have `cbegin()` and `cend()`, and functions like `erase` took an `iterator` instead of a `const_iterator`. This made `const_iterators` much less useful than they are now, and the (now obsolete) advice in *Effective STL* (2001, Item 26) was to avoid them.

Writing a vector using an iterator

We can rewrite our function to use sequential access via an iterator instead of random access:

```
// write vector to an output stream
void write_vector(ostream &out, const vector<double> &v) {
    out << "vector:";
    for (auto it = v.cbegin(); it != v.cend(); ++it)
        out << ' ' << *it;
    out << '\n';
}
```

Deleting zeroes using an iterator

```
// remove zero values from a vector
void delete_zeroes(vector<double> &v) {
    auto it = v.cbegin();
    while (it != v.cend())
        if (*it == 0)
            it = v.erase(it);
        else
            ++it;
}
```

Caution

We need to update `it` with the iterator returned by `erase`, because `erase` invalidates `it` and all later iterators.

This still has the same performance problem, but now it is easier to switch containers.

Standard sequential containers

vector implemented as an extensible array (like Java's `ArrayList`): efficiently supports **size**, **push_back**, **back**, **pop_back** and indexing.


deque implemented as an extensible array with wraparound (like Java's `ArrayDeque`): efficiently supports **size**, **push_back**, **back**, **pop_back**, **push_front**, **front**, **pop_front** and indexing.

list implemented as a doubly-linked list (like Java's `LinkedList`): efficiently supports **size**, **push_back**, **back**, **pop_back**, **push_front**, **front**, **pop_front**, **erase** and **insert**, but not indexing.

Unlike in Java, there is no subtyping involved: the different containers just use the same names for their member functions, which makes it easy to switch containers.

Implementation of `list`

The standard `list` type is a doubly linked list (see session 5 of Data Structures and Algorithms), e.g.



This structure supports inserting and removing elements at any point in constant time. However, it cannot support indexing efficiently.

Writing a list using an iterator

To adapt `write_vector` to `list`, we just change `vector` to `list`:

```
// write list to an output stream
void write_list(ostream &out, const list<double> &v) {
    out << "list:";
    for (auto it = v.cbegin(); it != v.cend(); ++it)
        out << ' ' << *it;
    out << '\n';
}
```

This works because all the operations used are common to the two containers.

Deleting all the zeroes from a list

Similarly, we can adapt `delete_zeroes` to lists:

```
// remove zero values from a list
void delete_zeroes(list<double> &v) {
    auto it = v.cbegin();
    while (it != v.cend())
        if (*it == 0)
            it = v.erase(it);
        else
            ++it;
}
```

However, because we are now using lists, `erase` is fast, and `delete_zeroes` takes time $O(n)$.

Range-based `for` loops

Shorthand: range-based `for`

Looping through the whole container is common, C++11 introduced an abbreviation for it:

- not changing the elements (for small types):

```
for (auto x : v)
    // do something with x
```

- not changing the elements, but avoiding a copy:

```
for (const auto &x : v)
    // do something with x
```

- changing the elements:

```
for (auto &x : v)
    x += 3;
```

Copy of the value of an element

A range-based **for** loop with a value variable

```
for (auto x : v) {
    // body looking at x
}
```

is equivalent to

```
for (auto it = v.cbegin(); it != v.cend(); ++it) {
    auto x = *it;
    // body looking at x
}
```

Use this when elements are small, and you don't want to change them.

Constant reference to an element

A range-based **for** loop with a **const** reference

```
for (const auto &x : v) {
    // body looking at x
}
```

is equivalent to

```
for (auto it = v.cbegin(); it != v.cend(); ++it) {
    const auto &x = *it;
    // body looking at x
}
```

Use this when elements are big, and you don't want to change them.

Reference to an element

A range-based `for` loop with a non-const reference

```
for (auto &x : v) {  
    // body updating x  
}
```

is equivalent to

```
for (it = v.begin(); it != v.end(); ++it) {  
    auto &x = *it;  
    // body updating x  
}
```

Use this when you want to change elements of the container.

Writing a vector using range-based `for`

The `write_list` function using a range-based `for`:

```
// write list to an output stream  
void write_list(ostream &out, const list<double> &v) {  
    out << "list:";  
    for (auto x : v)  
        out << ' ' << x;  
    out << '\n';  
}
```

Updating a vector using range-based `for`

If we want to modify the container:

```
// add one to each element  
void increment_list(list<double> &v) {  
    for (auto &x : v)  
        x += 1;  
}
```

The `&s` here avoid copying, but they are both also crucial to ensuring that elements of the container get updated:

- `v` is an alias for the list passed as an argument.
- `x` is an alias for an element of the list.

Associative containers

The `map` container

`map<K, V>` is an *associative container* (like Java's `TreeMap`). Maps can be efficiently indexed like vectors, but by type `K` (instead of `int`):


```
map<string, int> days;
days["January"] = 31;
days["February"] = 28;
days["March"] = 31;
...
string n = "October";
cout << n << " has " << days[n] << " days\n";
```

Caution

The expression `m[k]` creates an entry for `k` if none exists in `m` already.

Notional view of a map

A **map** stores an association of keys with values, in key order:

"April"	30
"August"	31
"December"	31
"February"	28
"January"	31
"July"	31
"June"	30
"March"	31
"May"	31
"November"	30
"October"	31
"September"	30

- Internally, **map** uses a form of balanced search tree (see sessions 7 and 8 of Data Structures and Algorithms).
- The **map** provides fast access by key.
- Iterators on the **map** provide sequential access to pairs of keys and values, in key order.

Counting occurrences of words

Task: read in words, and then print each unique word with the number of times it occurs in the input.

We shall use a map from words that we encounter to **int**:

```
map<string, int> count;
```

- This declares a map called **count**, and performs default initialization (as an empty map).
- In our program, the **int** will always be 1 or more – words that don't occur in the input won't have entries in the map.

Rest of the program:

1. read words, updating **count**
2. print contents of **count**

Computing the counts

The core of the the program is deceptively concise:

```
// read input words, updating their counts
string w;
while (cin >> w)
    ++count[w];
```

The loop is familiar, but there's a lot going on in the last line:

- If the map has an entry for **w**, then **count[w]** is a reference to it.
- If not, an entry for **w** is created, given default initialization (for **int**, set to 0), and then **count[w]** is a reference to the new entry.
- In either case, the entry is then increased by **++**.

At the end of this loop, **count[w]** is the number of times that **w** occurred in the input, for each **w** that occurred.

Printing the counts

It remains to go though the map, printing words and associated counts:

```
// write each word and its number of occurrences
for (const auto &p : count)
    cout << p.first << '\t' << p.second << '\n';
```

- The loop goes through the map in key order.
- At each step **it** refers to a **pair<const string, int>**, an object with members
 - **first**, a **const** of type **string**
 - **second** of type **int**
- We use **&**, to avoid copying the pair.
- We use **const**, to declare that we will not change the pair.

Iterator form

An equivalent form of that loop using iterators would be

```
for (auto it = count.cbegin(); it != count.cend(); ++it)
    cout << it->first << '\t' << it->second << '\n';
```

- We use **cbegin()** to get the **const_iterator**, because we are only looking at the values, not changing them.
- **it** has type **map<string, int>::const_iterator**.
- ***it** has type **pair<const string, int>**, with data members **first** and **second**.
- **it->first** is short for **(*it).first**

```
#include <iostream>
#include <string>
#include <map>

using namespace std;

int main() {
    // count of occurrences of each word in the input
    map<string, int> count;

    // read input words, updating their counts
    string w;
    while (cin >> w)
        ++count[w];

    // write each word and its number of occurrences
    for (const auto &p : count)
        cout << p.first << '\t' << p.second << '\n';
    return 0;
}
```

Figure 4.1: Program to count occurrences of words in the input

Exercises

1. Modify the **average** function from last week to use iterators instead of indices, and then adapt it to lists. Do this both as an iterator loop and as a ranged **for** loop.
2. Write a function

```
void scale(double s, list<double> &vs)
```

that modifies **vs**, multiplying each element by **s**. Again, do this both as an iterator loop and as a ranged **for** loop. (This will need **iterator** rather than **const_iterator**.)

3. Modify the **maximum** function from last week to use iterators instead on indices, and then adapt it to lists. (This will require a slightly different loop, because this function treats the first element specially.)
4. Write a function

```
void delete_first_zero(list<double> &vs)
```

that only deletes the first zero in the list, if any.

5. Write a function

```
void remdups(list<string> &ws)
```

that removes adjacent duplicate values from the list.

6. Modify the occurrence-counting program to print a histogram, i.e. next to each word, a row of n stars instead of a number n .
7. Write a program that reads in a series of word-number pairs, and outputs the median and average of the values paired with each number. For example, given the input

```
abc 3
def 4
abc 1
abc 4
```

it should produce the output

```
abc: average = 2.66667, median = 3
def: average = 4, median = 4
```

You will need to build a vector for each word, and then use the functions from last week to compute the average and median.